

JAYANTH DASAMANTHARAO

New Brunswick, NJ — (848)313-5617 — jayanthdasamantharao@gmail.com — [linkedin.com/in/djayanth](https://www.linkedin.com/in/djayanth)

PROFESSIONAL EXPERIENCE

AI Engineer - Fedway Associates, EliteUS Software Solutions Aug 2024 - Present | Edison, NJ

- Built large-scale relevance and ranking models using Python, SQL, ML techniques to personalize user-facing AI experiences, processing 5M+ interaction records and improving task relevance and 30% user engagement.
- Developed multimodal retrieval and understanding pipelines using transformer-based NLP and vision models to support semantic search, content understanding, cross-modal relevance in production systems.
- Applied large language models (GPT-4, LLaMA, Gemini) and RAG to improve reasoning quality, semantic understanding, response accuracy achieving 30%+ gains through iterative experimentation and evaluation.
- Deployed scalable ML infrastructure supporting training, evaluation and real-time inference using Python, FastAPI, cloud services, CI workflows, enabling rapid experimentation and reliable production deployment.

Data Science Intern, Harvest Software Solutions June 2023 - Sept 2023 | Jacksonville, Florida

- Engineered scalable ML pipelines and monitoring systems for large advertising datasets, increasing campaign visibility from 15% to 45% while ensuring reliability and performance.
- Built ML services and NLP APIs using XGBoost, H2O AutoML, and FastAPI, improving sentiment classification accuracy by 65% and campaign spend efficiency by 50%.
- Designed SQL/Python data pipelines for feature engineering, dataset validation, and model training, enabling repeatable experimentation and reliable deployment workflows.

Application Development Associate, Accenture June 2021 - Aug 2022 | Hyderabad, India

- Developed and optimized Python-based ETL and data processing pipelines using hash maps, sets, and indexed lookups to efficiently clean, deduplicate, transform large datasets, reducing 50% production issues.
- Implemented SQL workflows and joins with partitioning, indexing, and window functions, enabling efficient aggregation and analysis over millions of records in AWS Athena and Redshift.
- Applied batching, caching, and algorithmic optimizations for data validation and pipeline execution, improving system performance and reliability by 25%.

TECHNICAL SKILLS

Programming: Python, R, C++, SQL. **Frameworks:** PyTorch, TensorFlow, Hugging Face Transformers, FastAPI.
Machine Learning: End-to-end ML pipelines, model training, evaluation, monitoring, feature engineering, MLOps.
GenAI & LLMs: Transformer-based models, LLM fine-tuning, retrieval-augmented generation, multimodal modeling.
Cloud & MLOps: PostgreSQL, Data Warehousing, AWS (Bedrock, Athena, Redshift), GCP, Azure, Pinecone, Chroma.
Deep Learning: Neural Networks, CNNs, LSTMs, YOLO, Faster R-CNN, GANs, Time-Series Forecasting.

EDUCATION

Rutgers University, Master of Science in Data Science Sept 2022 - May 2024 | New Brunswick, NJ
Personal Website: <https://jayanthdasamantharao.github.io/>

Andhra University B.Tech in Electrical and Electronics Engineering 2017 - 2021 | Visakhapatnam, India

RESEARCH PROJECTS

Self-Driving Cars (Computer Vision): Built autonomous driving cars with vision pipeline using CNNs, YOLO, Faster R-CNN, achieving 95%+ object detection accuracy.

Multimodal Retrieval & Question Answering System (Production): Built LLM-assisted retrieval and question answering pipelines using embeddings, multimodal inputs, and evaluation-driven iteration to improve response quality and factual grounding in production systems.

Bayesian Classification: Implemented Bayesian logistic regression in PyMC3 to predict diabetes status, incorporating informative priors, achieving 0.85 F1-score.

For additional projects and code samples: <https://github.com/JayanthDasamantharao>